Gemini Enterprise
+ Sensitive Data Protection (Preview)
User Guide

Shared under NDA

Private Preview

This feature is subject to the "Pre-GA Offerings Terms" in the General Service Terms section of the Service Specific Terms. Pre-GA features are available "as is" and might have limited support. For more information, see the launch stage descriptions.

- Overview
  - The Problem: The Security gap in Enterprise Data
  - The Solution: Closing the loop with Inbound Content Policies
  - Key features of Content Policy
  - Content Policy In Action
- Getting Started
  - Step 1 & 2
  - Step 3: Allow Listing
    - Allowlisting
  - Step 4: Creating and Using a Content Policy
    - Permissions
    - Create a Content Policy
    - (Optional) Update a Content Policy
    - Detecting Metadata Label
  - Step 5: Assigning and Using a Content Policy with Gemini Enterprise
    - Permissions
    - Associate Content Policy with a Data Connector
    - Associate Content Policy with Notebook LM
- Known Issues

Overview

The Problem: The Security gap in Enterprise Data

Connecting Gemini to your enterprise data—like SharePoint, Jira, and Google Workspace—is what makes it a powerhouse for your team. However, most security strategies currently focus only on the "front door": the prompts users type and the responses the AI gives back.

The problem is that even if a user asks a safe question, the AI might pull in "inbound" data from your connectors that shouldn't be processed—like documents with sensitive internal labels, inappropriate content, or unintended PII. Relying solely on prompt scanning means you're missing the chance to govern the data **at the source** before the AI even interacts with it. To truly support your employees, you need a way to ensure the data being "fed" to the AI is as secure as the conversation itself.

The Solution: Closing the loop with Inbound Content Policies

This preview introduces a second, essential layer of protection: **Inbound Data Content Policies.** Instead of just watching the chat, you can now place a "guard" on your data connectors.

This creates a two-pathway defense:

- **Pathway 1 (The New Layer):** A policy for "inbound" data. This automatically checks the content and metadata (like Microsoft Sensitivity Labels) of files fetched from your connectors.
- **Pathway 2 (The Standard Layer):** A policy for the user interaction itself, scanning the prompts and responses in real-time.

By combining both, you ensure that the data the AI uses is just as vetted as the answers it provides. For information on that pathway see here.

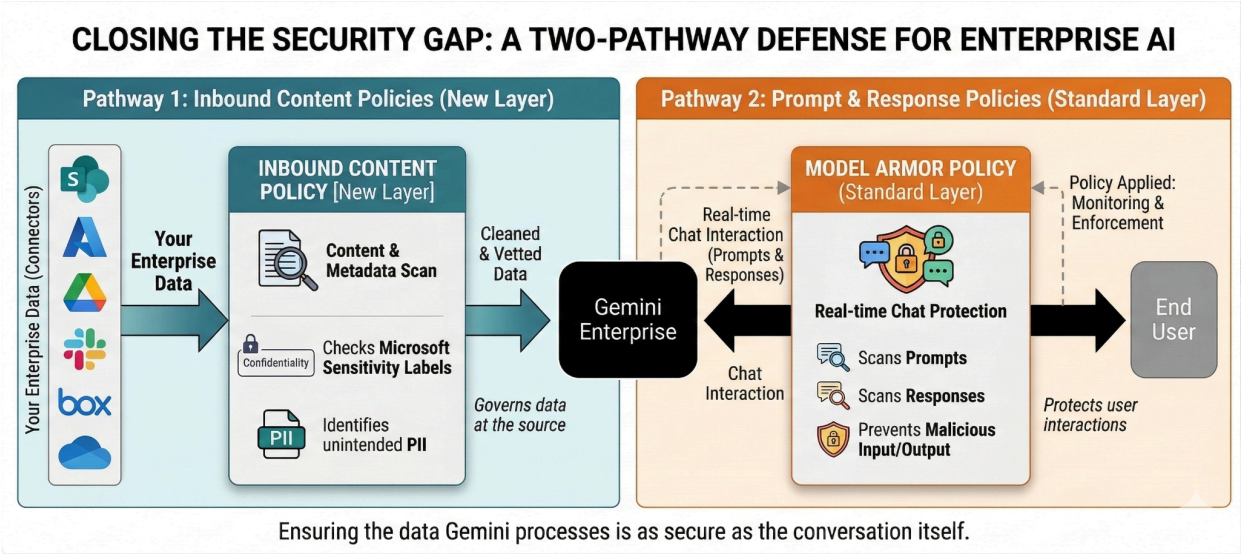

Key features of Content Policy

Sensitive Data Protection Content Policies bring you the power of Sensitive Data Protection content inspection along with ability to create allow and block rules. This way you can, for example, inspect data for violating or unwanted content and choose when and how to allow or block this content.

Key features:

- Support for images including:
    - Analysis of textual content using AI/ML powered Optical Character Recognition (OCR)
    - Image object detention for non textual objects like people, faces, ID cards, and more
    - Image Safety detection for explicitly, racy, or violent imagery

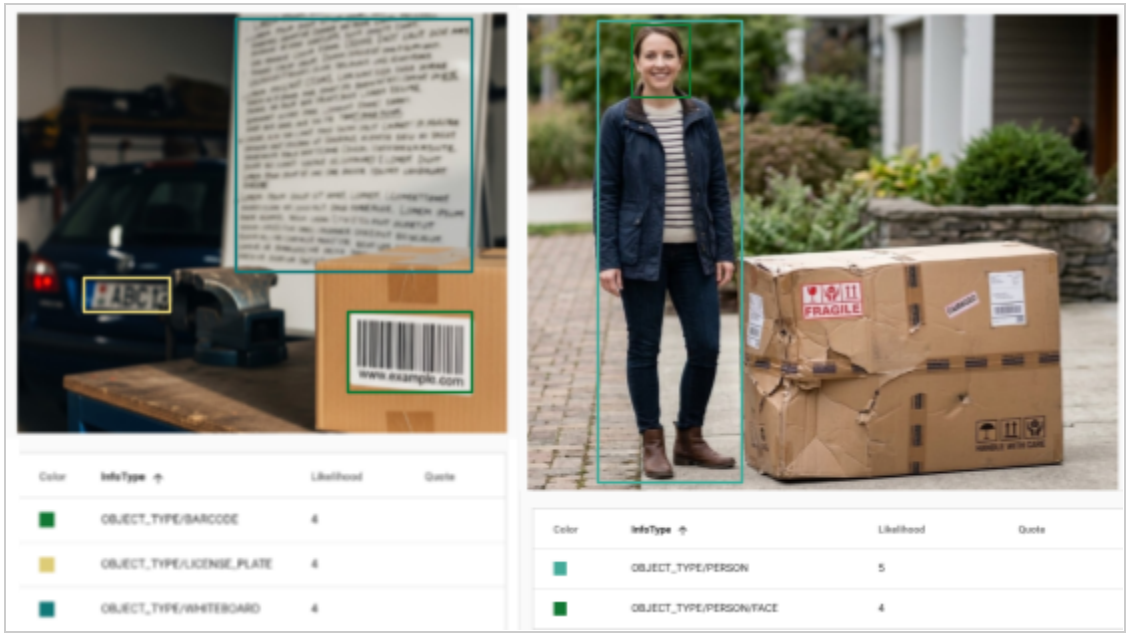

- Support for rich documents (PDF, DOCx, PPTx, XLSx)
    - Analysis of textual content using extracted from PDFs
    - Support for embedded images
    - Support for metadata data labels like MIP/AIP Labels

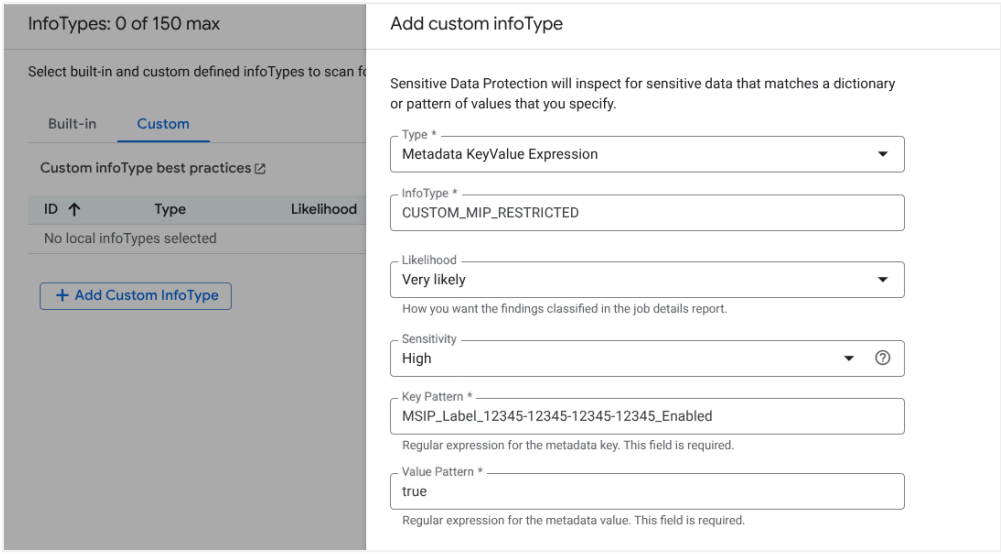

- Support for unstructured and structured text analysis with over [200 built-in infoTypes]() including new AI/ML context detectors, source code detection, and more

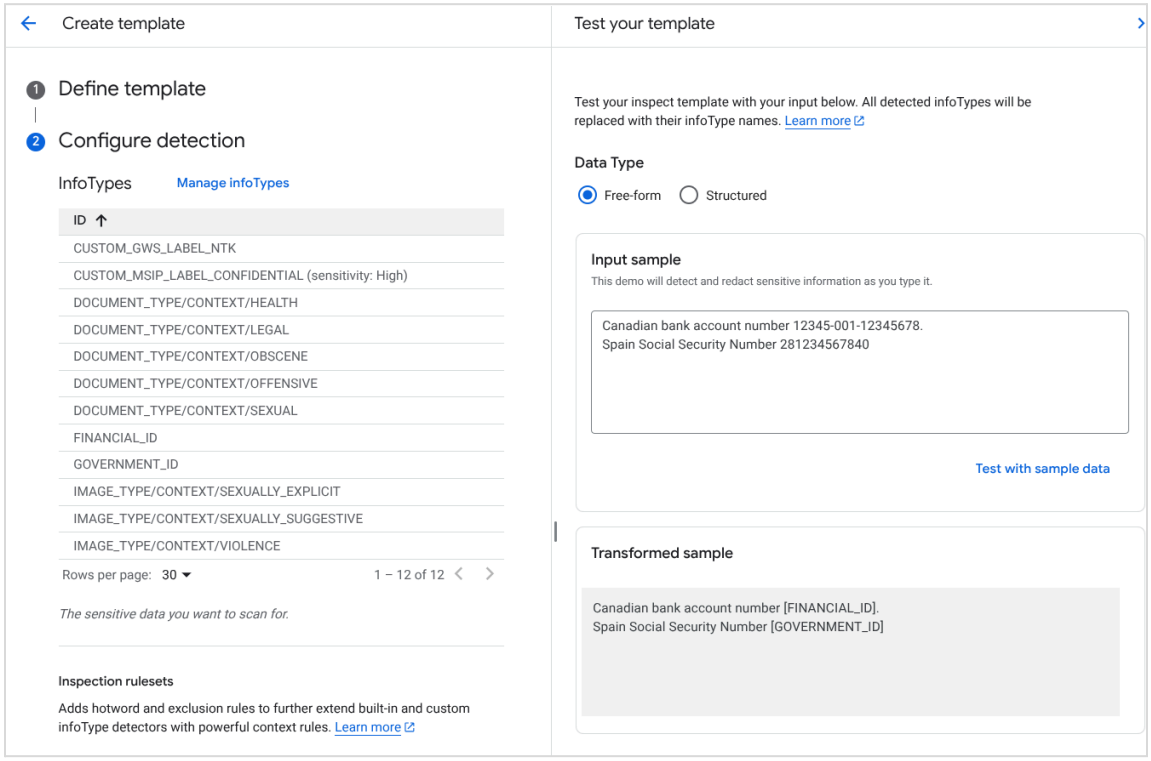

- Support for custom infoTypes (e.g. regex, dictionary, and more)
- Support for custom Rules (exclusion, adjustment, hotword)

Content Policy In Action

Once a content policy is applied, content that violates your configured policy will be blocked.

For example, below is an example interaction with Gemini Enterprise where a request was blocked because the data coming from the connector violated policy. This illustrates the enhanced security of having policy on pathways 1 and 2.

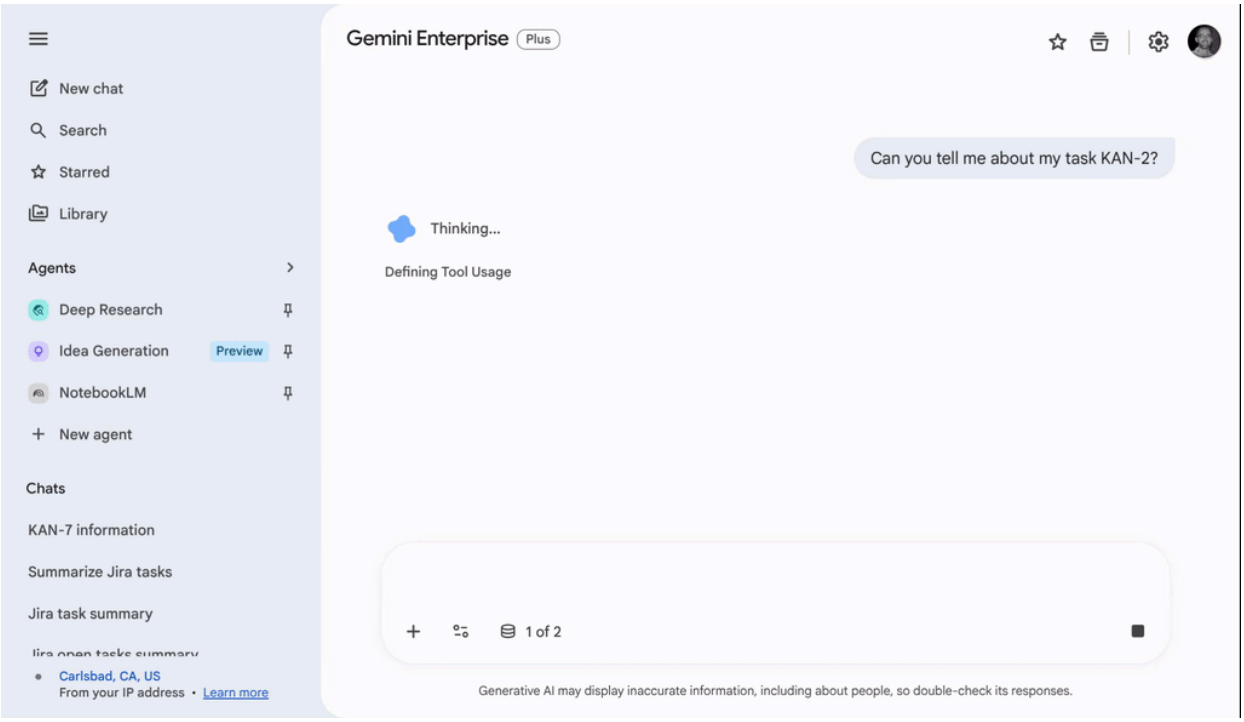

Without pathway 1, Gemini Enterprise would still have access to the restricted content and could leverage it in summaries or responses. In the example above, the ticket for KAN-2 contains sensitive customer data like payment card information.

In NotebookLM content such as uploads that violate policy will appear like this indicating that they were not ingested into your notebook:

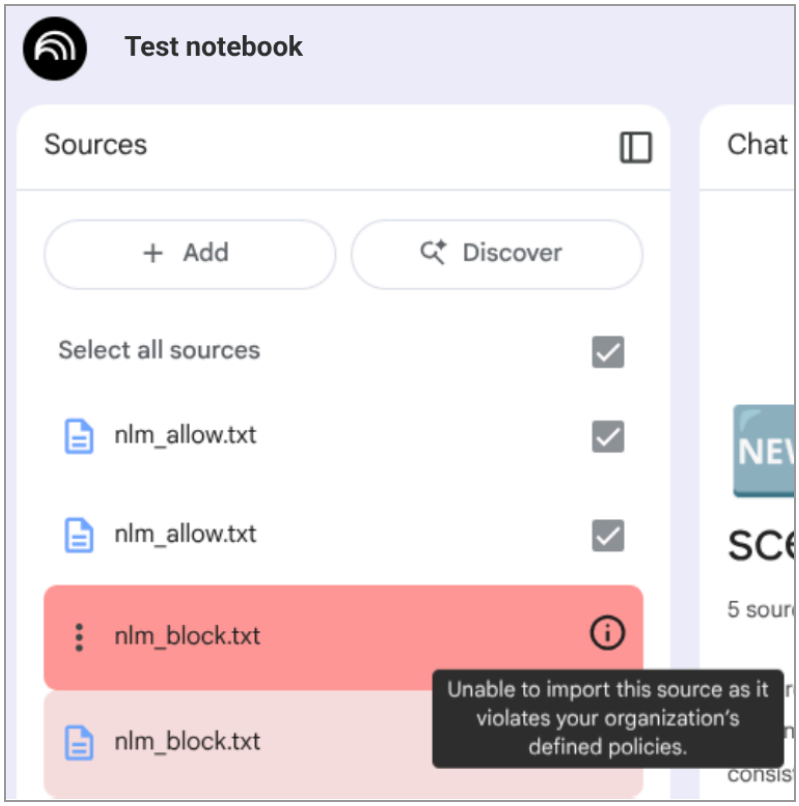

Getting Started

To configure Sensitive Data Protection policy and start using it in Gemini Enterprise, the following steps are needed:

- **Step 1:** You have Gemini Enterprise set up and running
- **Step 2:** You have at least one data connector linked to your Gemini Enterprise App or Notebook LM configured
- **Step 3:** Your account team requests for your projects where Gemini Enterprise and connectors are configured to be on the preview allowlist
- **Step 4:** Create a Sensitive Data Protection "Content Policy" using Console UI and API
- **Step 5:** Assign this Content Policy to:
    - Data connectors (using the API)
    - NotebookLM instance (using the API)

Once these steps are completed, the following user interactions will be governed by Sensitive Data Protection Content Policy:

- NotebookLM: file uploads
- Gemini Enterprise federated connector: search, assist, retrieval.

Below are detailed instructions and examples for creating content policy and assigning it to your connectors and Notebook LM instances.

Step 1 & 2

See Documentation on Gemini Enterprise here:
https://docs.cloud.google.com/gemini/enterprise/docs

Step 3: Allow Listing

Allowlisting

To use this feature you will need to be on the allow list for the preview. Please contact your account team to join.

Support Group: agentspace-sdp-support@google.com

Step 4: Creating and Using a Content Policy

A content policy operates in a manner comparable to the existing content.inspect and hybrid inspection methodologies, providing synchronous results upon the submission of content. While the inspection rules and configurations remain consistent with those defined in an inspectConfig, the primary distinction lies in the output: a content policy request generates a formal verdict based on the presence of specific findings. In contrast, traditional content and hybrid methods produce a comprehensive list of identified findings and not a final verdict.

Permissions

To create or update content policies:

- Required permissions:
  - dlp.contentPolicies.create
  - dlp.contentPolicies.update
- Included in roles:
  - dlp/admin
  - dlp/contentPoliciesEditor

To apply a content policy:

- Required permissions:
  - dlp.contentPolicies.apply
- Included in roles:
  - dlp/user
  - dlp/contentPoliciesConsumer

- Higher-privilege roles like dlp/contentPoliciesEditor or dlp/admin

Create a Content Policy

When you create a content policy it will exist as a resource. Once a content policy is created, you can make a request against that policy to get a verdict. For this preview those requests will only be made as part of Gemini Enterprise integration. . Policies can later be updated or deleted.

Structure of a content policy creation request:

- **parent**: the project and location where the content policy will be created.
- **content_policy_id** : a unique identifier for the policy. If one is not provided a randomly generated ID will be created and used.
- **inspect_template**: the body of an inspect template including an inspect config. Unlike other methods, this config is inlined for efficiency.
  Note: if you do not have an inspection template or want to create a new one, you can do this using Terraform, via API, or through Cloud Console here. Once you have this created the inspect template you can download this
- **rules**: conditions based on infoTypes findings and a verdict. Supported verdicts are ALLOW or BLOCK

The policy below instructs Sensitive Data Protection to look for credit card numbers only (using infoType **CREDIT_CARD_NUMBER**) and to block if anything is found.

```shell
PROJECT_ID="your-project-id"
LOCATION="global" # or 'us', 'eu'
SDP_POLICY="your_sdp_content_policy_full_resource_name"

# Get your access token
ACCESS_TOKEN=$(gcloud auth print-access-token)

# The PATCH request
curl -X POST \ -H "Authorization: Bearer ${ACCESS_TOKEN}" \
-H "Content-Type: application/json" \
-H "x-goog-user-project:<your-project-id>" \
"https://dlp.googleapis.com/v2/projects/${PROJECT_ID}/locations/${LOCATION}/contentPolicies" \
-d '{  "content_policy": {    "inspect_template": {     "display_name": "Sample Inspect Configuration
for Content Policy",    "inspect_config": {       "info_types": [         {           "name":
```

```
"CREDIT_CARD_NUMBER"        }      ],      "min_likelihood": "POSSIBLE"     }   },   "rules": {
"conditions": {      "info_type_condition": {        "any_info_type": {}      }     },
    "return_verdict": "BLOCK"    } },  "content_policy_id": "${SDP_POLICY}"}'
```

Example Formatted JSON Request Payload with policy id “test_policy1”

```JSON
{
  "content_policy": {
    "inspect_template": {
      "display_name": "Sample Inspect Configuration for Content Policy",
      "inspect_config": {
        "info_types": [
          {
            "name": "CREDIT_CARD_NUMBER"
          }
        ],
        "min_likelihood": "POSSIBLE"
      }
    },
    "rules": {
      "conditions": {
        "info_type_condition": {
          "any_info_type": {}
        }
      },
      "return_verdict": "BLOCK"
    }
  },
  "content_policy_id": "test_policy1"
}
```

Here is an example response to the create request:

```
{
  "name": "projects/<your-project-id>/locations/us/contentPolicies/test_policy1",
  "inspectTemplate": {
    "displayName": "Sample Inspect Configuration for Content Policy",
    "inspectConfig": {
      "infoTypes": [
        {
```

```
        "name": "CREDIT_CARD_NUMBER"
      }
    ],
    "minLikelihood": "POSSIBLE",
    "limits": {}
  }
 },
 "rules": [
  {
    "conditions": [
      {
        "infoTypeCondition": {
          "anyInfoType": {}
        }
      }
    ],
    "returnVerdict": "BLOCK"
  }
 ]
}
```

(Optional) Update a Content Policy

In this example below we update/patch the policy:

```shell
PROJECT_ID="your-project-id"
LOCATION="global" # or 'us', 'eu'
SDP_POLICY="your_sdp_content_policy_full_resource_name"

# Get your access token
ACCESS_TOKEN=$(gcloud auth print-access-token)

# The PATCH request
curl -X PATCH \ -H "Authorization: Bearer ${ACCESS_TOKEN}" \
-H "Content-Type: application/json" \
-H "x-goog-user-project:<your-project-id>" \
"https://dlp.googleapis.com/v2/projects/${PROJECT_ID}/locations/${LOCATION}/contentPolicies/${SDP_POLICY}" \
-d '{  "content_policy": {    "inspect_template": {     "display_name": "Updated Inspect Configuration for Content Policy",     "inspect_config": {       "info_types": [         {           "name": "CREDIT_CARD_NUMBER"         }       ],       "min_likelihood": "LIKELY"     }   },   "rules": {     "conditions": {       "info_type_condition": {         "any_info_type": {}       }     },     "return_verdict": "BLOCK"   } }}'
```

Example JSON Request Payload

```JSON
{
  "content_policy": {
    "inspect_template": {
      "display_name": "Updated Inspect Configuration for Content Policy",
      "inspect_config": {
        "info_types": [
          {
            "name": "CREDIT_CARD_NUMBER"
          }
        ],
        "min_likelihood": "LIKELY"
      }
    },
    "rules": {
      "conditions": {
        "info_type_condition": {
          "any_info_type": {}
        }
      },
      "return_verdict": "BLOCK"
    }
  }
}
```

Detecting Metadata Label

To detect a metadata label like a MIP/AIP/MSIP label. (Microsoft Sensitive Label), you would leverage a new custom infoType of type = “metadata_key_value_expression” as part of your inspectConfig. Below is a sample request snippet showing this new custom infoType.

```JSON
...
inspect_config {
  custom_info_types {
    info_type {
      name: "CUSTOM_MSIP1"
    }
    likelihood: VERY_LIKELY
    metadata_key_value_expression {
```

```
      key_regex: "MSIP_Label_12345-12345-1234-12345_Enabled"
      value_regex: "true"
    }
  }
}
...
```

This will look for the MIP/AIP/MSIP label with GUID = “12345-12345-1234-12345” being found and enabled in a document. The result will be a finding of the type “CUSTOM_MSIP1” as defined.

Note, infoTypes of this type are not currently supported in inspectionRuleSets but can be used in ContentPolicy rules.

Step 5: Assigning and Using a Content Policy with Gemini Enterprise

Next you need to assign or associate your Content Policy with your Gemini Enterprise environment.

Permissions

Customers’ admin need to grant Gemini Enterprise service account (usually in the format of ***service-{PROJECT_ID}@gcp-sa-discoveryengine.iam.gserviceaccount.com***) with one of the following roles:

1) **dlp/user**
2) **dlp/admin**

Please note that by default we also support cross-project reference if they are in the same vpc-sc perimeter. The Sensitive Data Protection content policy created under projectA can be used for Gemini Enterprise in projectB, as long as in projectA, you grant the one of the above roles to **service-{projectB}@gcp-sa-discoveryengine.iam.gserviceaccount.com**

Associate Content Policy with a Data Connector

The following command will associate a Content Policy with an already existing Data Connector

```
Shell
```

```
PROJECT_ID="your-project-id"
LOCATION="global" # or 'us', 'eu'
COLLECTION_ID="your_connector_collection"
SDP_POLICY="your_sdp_content_policy_full_resource_name"

# Get your access token
ACCESS_TOKEN=$(gcloud auth print-access-token)

# The PATCH request
curl -X PATCH \ -H "Authorization: Bearer ${ACCESS_TOKEN}" \ -H "Content-Type:
application/json" \
"https://discoveryengine.googleapis.com/v1alpha/projects/${PROJECT_ID}/locations/${LOCATIO
N}/collections/${COLLECTION_ID}/dataConnector?updateMask=dataProtectionPolicy" \
-d '{ "dataProtectionPolicy": { "sensitiveDataProtectionPolicy": { "policy": "${SDP_POLICY}" } } }'
```

Please note that the Sensitive Data Protection policy region needs to match with the Gemini Enterprise region.

Associate Content Policy with Notebook LM

To enable Sensitive Data Protection in NotebookLM, use the curl commands below. Note where these commands reference the project "number" vs the project "id":

```Shell
# Replace placeholders with your actual values
export PROJECT_NUM="YOUR_PROJECT_NUM" # MODIFY THIS
export PROJECT_ID="YOUR PROJECT ID" # MODIFY THIS
export LOCATION="YOUR LOCATION" # Only supports "global", "us", "eu"
# Format 'projects/{project}/locations/{location}/contentPolicies/{policy_id}'
export USER_POLICY="YOUR user policy"


curl -X PATCH \
-H "Authorization: Bearer $(gcloud auth print-access-token)" \
-H "Content-Type: application/json" \
-H "x-goog-user-project: ${PROJECT_ID}" \
"https://${LOCATION}-discoveryengine.googleapis.com/v1alpha/projects/${PROJECT_NUM}?upd
ate_mask=customer_provided_config" \
-d '{
"customer_provided_config": {
  "notebooklm_config": {
```

```
    "data_protection_policy": {
      "sensitive_data_protection_policy": {
        "policy": "${USER_POLICY}",
      }
    }
  }
}
}'
```

Example JSON Request Payload

```JSON
{
  "customer_provided_config": {
    "notebooklm_config": {
      "data_protection_policy": {
        "sensitive_data_protection_policy": {
"policy":"projects/<your-project_id>/locations/global/contentPolicies/test_policy1"
      }
      }
    }
  }
}
```

Here is an example response to the create request:

```
{
  "name": "projects/<your project number>",
  "createTime": "2025-03-14T17:01:10.660901312Z",
  "provisionCompletionTime": "2025-03-14T17:01:10.660901312Z",
  "serviceTermsMap": {
    "GA_DATA_USE_TERMS": {
      "id": "GA_DATA_USE_TERMS",
      "version": "2022-11-23",
      "state": "TERMS_ACCEPTED",
      "acceptTime": "2025-03-14T17:01:10.660901312Z"
    }
  },
  "customerProvidedConfig": {
    "notebooklmConfig": {
      "dataProtectionPolicy": {
        "sensitiveDataProtectionPolicy": {
          "policy": "projects/dlp-onboarding-test1/locations/global/contentPolicies/test_policy1"
        }
      }
    }
```

```
  }
}
```

Please wait ~ 10 minutes for the configuration change to propagate. NotebookLM will now start using the specified Sensitive Data Protection policies to scan the uploaded source documents.

Details on Supported Pathways

Features	Federated Connectors	Indexed Connectors	NLM Uploads	GE “talk to content” uploads
Image OCR (text extraction and analysis)	Yes	Coming Soon	Yes	Coming Soon
Image Object Detection (e.g. person, barcode, id card)	Yes	Coming Soon	Yes	Coming Soon
Image Safety	Yes	Coming Soon	Yes	Coming Soon
Rich Document Support (including embedded images)	Yes	Coming Soon	Yes	Coming Soon
AIP / MIP labels support for PDF, DOCx, XLSx, PPTx	Yes	Coming Soon	Yes	Coming Soon
Text Content Analysis	Yes	Coming Soon	Yes	Via Model Armor
Policy Type	*Sensitive Data Protection* **Content Policy**	*Sensitive Data Protection* **Content Policy**	*Sensitive Data Protection* **Content Policy**	Model Armor

Known Issues

Indexed Connectors are not currently supported but coming soon. Until then you may receive an error

```
{
  "error": {
    "code": 400,
    "message": "Data protection policy is only supported for federated connectors.",
    "status": "INVALID_ARGUMENT"
  }
```

}

Uploading files directly into Gemini Enterprise from the “+” button are currently only scanned through Model Armor and have the following limitations:
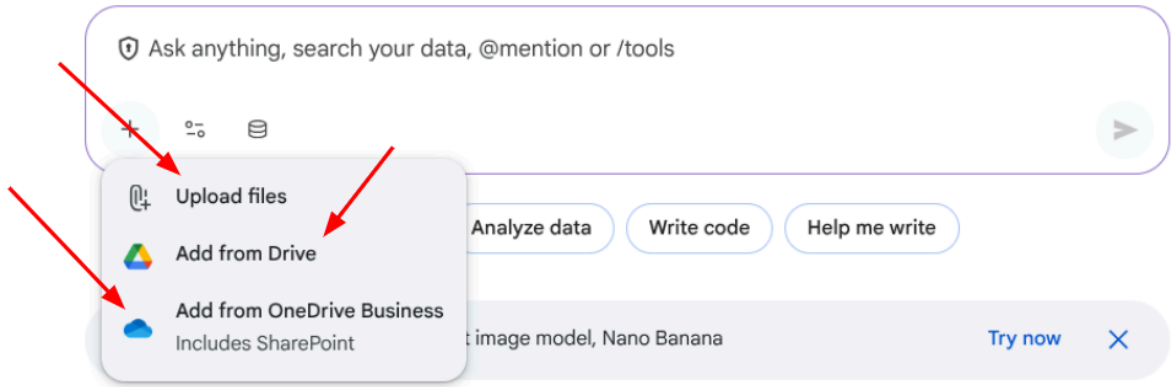
- Images are not supported
- Rich documents (PDF, Docs, etc.)
  - Images in documents are not supported
  - Metadata Labels (eg. MIP/AIP labels) are not supported